

1 PyMC-Marketing: Bayesian Marketing Mix Models 2 and Customer Analytics in Python

3 William Dean  ¹, Juan Orduz  ¹, Colt Allen¹, Carlos Trujillo¹, Ricardo
4 Vieira¹, Benjamin T. Vincent¹, and Thomas Wiecki¹

5 1 PyMC Labs

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- 6 ▪ [Review](#) ↗
- 7 ▪ [Repository](#) ↗
- 8 ▪ [Archive](#) ↗

9 **Editor:** [Open Journals](#) ↗

10 **Reviewers:**

- 11 ▪ [@openjournals](#)

12 **Submitted:** 01 January 1970

13 **Published:** unpublished

License

14 Authors of papers retain copyright
15 and release the work under a
16 Creative Commons Attribution 4.0
17 International License ([CC BY 4.0](#))
18

19 Statement of Need

Marketing organizations struggle to attribute sales outcomes to specific marketing activities across multiple touchpoints and delayed conversion effects. Existing solutions suffer from: (1) black-box proprietary models with limited customization; (2) oversimplified approaches failing to capture marketing dynamics; and (3) lack of uncertainty quantification for high-stakes decisions.

PyMC-Marketing addresses these gaps by bridging marketing science research and practical applications. It operationalizes advanced Bayesian methods—hierarchical modeling, experimental calibration, and uncertainty quantification—within a user-friendly, scikit-learn compatible API. Key innovations include time-varying coefficients using modern Gaussian process approximations optimized for marketing contexts, and a novel experimental calibration framework that integrates lift test results directly into model likelihood. While frequentist approaches like Robyn provide bootstrap-based intervals, all PyMC-Marketing outputs include full Bayesian posterior distributions, enabling decision-makers to assess risk explicitly.

29 Installation and Dependencies

30 PyMC-Marketing is available via conda-forge and pip. Core dependencies include PyMC (5.0),
31 NumPy, Pandas, ArviZ ([Kumar et al., 2019](#)), and scikit-learn. Optional dependencies enable
32 GPU acceleration (JAX), advanced samplers (NumPyro, Nutpie), and production deployment
33 (MLflow ([Zaharia et al., 2018](#)), Docker).

34 Key Features

35 PyMC-Marketing provides four distinct modules addressing comprehensive marketing analytics:

- 36 **1. Media Mix Modeling (MMM):** Multiple adstock functions, saturation curves, time-varying
37 parameters via HSGP ([Solin & Särkkä, 2020](#)), experimental calibration for causal inference,
38 budget optimization with business constraints, time-slice cross-validation, and marginal effects
39 analysis ([Arel-Bundock et al., 2024](#)).
- 40 **2. Customer Lifetime Value (CLV):** BTYD models ([Fader & Hardie, 2020](#)) including BG/NBD,
41 Pareto/NBD, and Gamma-Gamma frameworks with hierarchical extensions and individual-level
42 uncertainty.
- 43 **3. Bass Diffusion Models:** Product adoption forecasting ([Bass, 1969](#)) with flexible parameteri-
44 zation for innovation and imitation effects across multiple products.
- 45 **4. Customer Choice Models:** Discrete choice analysis ([Train, 2009](#)) based on random utility
46 theory, including multinomial logit and multivariate interrupted time series models.
- 47 **Production Ready:** All modules feature MLflow ([Zaharia et al., 2018](#)) integration, Docker
48 containerization, multiple MCMC backends (NumPyro ([Bingham et al., 2019](#)), Nutpie ([Seyboldt
49 & PyMC Developers, n.d.](#))), variational inference (ADVI), MAP estimation, data connectors
50 (e.g., Fivetran), and comprehensive diagnostics via ArviZ.

51 Key Advantages

52 PyMC-Marketing provides uncertainty quantification through full posterior distributions, exper-
53 imental calibration anchoring observational models to causal ground truth, and flexible budget
54 optimization with business constraints. The scikit-learn compatible API ensures seamless
55 integration into existing data science workflows. The library has been successfully deployed by
56 companies including HelloFresh and Bolt for production marketing analytics. Comprehensive
57 tutorials, example notebooks, and video resources are available in the online documenta-
58 tion at <https://www.pymc-marketing.io/en/stable/>, with community support from over 70
59 contributors and translations in Spanish.

60 Community Guidelines

- 61
 - **Issues:** Report bugs and feature requests on [GitHub](#)
 - **Discussions:** Join community discussions on [GitHub Discussions](#)
 - **Documentation:** Comprehensive guides and tutorials available at [pymc-marketing.io](#)
 - **Contributing:** See [CONTRIBUTING.md](#) for development guidelines
 - **Support:** Professional consulting available through [PyMC Labs](#)

66 Funding

67 PyMC-Marketing is a community-driven project built primarily through volunteer contributions
68 from over 70 developers. Some contributions by PyMC Labs affiliates have received partial
69 funding.

70 Acknowledgments

71 We acknowledge the PyMC development team for the foundational probabilistic programming
72 framework and the broader PyData ecosystem contributors. Special recognition goes to
73 the marketing science research community for developing the theoretical frameworks that
74 PyMC-Marketing operationalizes ([Bass, 1969](#); [Fader et al., 2005](#); [Jin et al., 2017](#); [Train, 2009](#)).

75 References

- 76 Arel-Bundock, V., Greifer, N., & Heiss, A. (2024). How to interpret statistical models
77 using marginaleffects for R and Python. *Journal of Statistical Software*, 111(9), 1–32.
78 <https://doi.org/10.18637/jss.v111.i09>
- 79 Bass, F. M. (1969). A new product growth for model consumer durables. *Management Science*,
80 15(5), 215–227. <https://doi.org/10.1287/mnsc.15.5.215>
- 81 Bingham, E., Chen, J. P., Jankowiak, M., Obermeyer, F., Pradhan, N., Karaletsos, T.,
82 Singh, R., Szerlip, P. A., Horsfall, P., & Goodman, N. D. (2019). Pyro: Deep universal
83 probabilistic programming. *Journal of Machine Learning Research*, 20(28), 1–6. <http://jmlr.org/papers/v20/18-403.html>
- 85 Fader, P. S., & Hardie, B. G. (2020). Customer lifetime value: What's next? *Journal of
86 Interactive Marketing*, 49, 26–37. <https://doi.org/10.1016/j.intmar.2019.08.001>
- 87 Fader, P. S., Hardie, B. G., & Lee, K. L. (2005). Counting your customers: Who are
88 they and what will they do next? *Management Science*, 51(9), 1371–1387. <https://doi.org/10.1287/mnsc.51.9.1371>
- 90 Google Inc. (2023). *Meridian: Bayesian media mix modeling at scale*. <https://github.com/google/meridian>. <https://github.com/google/meridian>
- 92 Jin, Y., Wang, Y., Sun, Y., Chan, D., & Koehler, J. (2017). *Bayesian methods for media
93 mix modeling with carryover and shape effects*. Google Inc. <https://research.google/pubs/pub46001/>
- 95 Kumar, R., Carroll, C., Hartikainen, A., & Martin, O. A. (2019). ArviZ a unified library for
96 exploratory analysis of bayesian models in python. *Journal of Open Source Software*, 4(33),
97 1143. <https://doi.org/10.21105/joss.01143>
- 98 Meta Platforms Inc. (2022). *Robyn: Marketing mix modeling*. <https://github.com/facebookexperimental/Robyn>. <https://github.com/facebookexperimental/Robyn>
- 100 Salvatier, J., Wiecki, T. V., & Fonnesbeck, C. (2016). Probabilistic programming in python
101 using PyMC3. *PeerJ Computer Science*, 2, e55. <https://doi.org/10.7717/peerj-cs.55>
- 102 Seyboldt, A., & PyMC Developers. (n.d.). *nutpie*. <https://github.com/pymc-devs/nutpie>
- 103 Solin, A., & Särkkä, S. (2020). Hilbert space methods for reduced-rank gaussian
104 process regression. *Statistics and Computing*, 30(2), 419–446. <https://doi.org/10.1007/s11222-019-09886-w>
- 106 Train, K. E. (2009). *Discrete choice methods with simulation*. Cambridge University Press.
107 <https://doi.org/10.1017/CBO9780511805271>
- 108 Zaharia, M. A., Chen, A., Davidson, A., Ghodsi, A., Hong, S. A., Konwinski, A., Murching,
109 S., Nykodym, T., Ogilvie, P., Parkhe, M., Xie, F., & Zumar, C. (2018). Accelerating
110 the Machine Learning Lifecycle with MLflow. *IEEE Data Eng. Bull.*, 41, 39–45. <https://api.semanticscholar.org/CorpusID:83459546>